

The Architecture of Artificial Learning: From McCulloch-Pitts Neurons to BURT-IMMA

Eighty years of iterative error correction—and a thirteen-layer architecture that forks the lineage: Matrix-Memory Equilibrium Propagation, a CIFG matrix-memory cell, CUDA kernels for a single RTX 3080, and a Lean 4 formal ledger

Jessica Westerhoff¹

¹Snap Kitty Collective · SnapKitty West, jessica@collectivekitty.com

2026 *Incorporating: Adversarial Mathematical Audit responses and corrected proof ledger
(August 2026)*

Evidence or Silence. Four epistemic labels appear throughout: **PROVED** (Lean 4 proof, zero **sorry**); **VERIFIED** (executable test, passes); **SKETCH** (formal statement, **sorry** on its path); **HYPOTHESIS** (empirical, falsification criterion specified). CUDA performance numbers are engineering characterization on RTX 3080 hardware, not audited benchmarks.

Abstract

Modern artificial intelligence rests on a single foundational mechanic: iterative error correction across distributed weights. **Part I** synthesizes that evolution—McCulloch-Pitts binary neurons, Rosenblatt’s trainable Perceptron, the linear-separability bottleneck, backpropagation, AlexNet’s convolutional breakthrough, the LSTM, self-attention, and feature superposition. **Part II** presents the fork.

BURT-IMMA is a thirteen-layer architecture whose learning rule is **Matrix-Memory Equilibrium Propagation (MMEP)**: a two-phase, locally computable, Hebbian-style update derived from the difference between a free equilibrium and a nudged equilibrium. No gradient tape. No weight transport. The answer is already in the fixed point.

At its heart sits a direct LSTM descendant—a Coupled Input-Forget Gate (CIFG) cell whose state is a full $d \times d$ matrix:

$$C_t = f_t \odot C_{t-1} + (1 - f_t) \odot (v_t \otimes k_t).$$

Every component is documented: the CIFG matrix memory and sum-inversion retrieval bounds; the MMEP energy function and convergence argument; the entropy-constrained softmax holding router entropy at $H(\alpha) \leq 0.20$ nats; the spectral projection that keeps relaxation contractive; the SmoothLeaky activation and its four axioms; Boolean perceptron actors under Huntington’s postulates; superpositioned induction heads; seven custom CUDA kernels for `sm_86`; and the Lean 4 formalization.

Audit note (August 2026). An adversarial mathematical audit applied multi-agent debate and first-principles re-derivation to the equations, theorems, and empirical claims. Six findings were produced. This edition integrates those findings, presents corrected claims, and documents which findings have been patched by subsequent Lean 4 machine-checked proofs. The corrected proof ledger appears in Section 12.

Contents

I	The Lineage: Eighty Years of Nudging Weights	3
1	The Foundational Mechanics of Artificial Neurons	3
1.1	The McCulloch-Pitts Abstraction (1943)	3
1.2	Rosenblatt and the Perceptron (1957)	3
1.3	The Linear-Separability Bottleneck (1969)	3
2	Multi-Layer Perceptrons and the Credit Assignment Problem	3
2.1	The Problem	3
2.2	Backpropagation and Gradient Descent	4
3	Evolution of Specialized Architectures	4
3.1	Convolutional Networks and AlexNet (2012)	4
3.2	Recurrent Networks, Vanishing Gradients, and the LSTM	4
4	The Transformer Era	4
4.1	Self-Attention: Query, Key, Value	4
4.2	Feature Superposition and Polysemantic Neurons	4
II	The Fork: BURT-IMMA and Matrix-Memory Equilibrium Propagation	5
5	Where and Why the Lineage Forks	5
5.1	The Thirteen Layers	5
6	The Memory Cell: The LSTM, Lifted	5
6.1	The IMMA Expert Cell	6
6.2	Sum-Inversion: Storage and Retrieval Theory	6
6.3	Trace Conservation (Corrected per Audit Finding 1)	6
7	MMEP: Credit Assignment Without the Tape	7
7.1	Equilibrium Propagation, Inherited and Extended	7
7.2	Energy Function	7
7.3	The Two Phases and the Whole Rule	7
7.4	Convergence (Audit-Strengthened)	7
7.5	Four-Phase Training Protocol	8
8	The Constraint Subsystem	8
8.1	$H(\alpha) \leq 0.20$ Nats, Always	8
8.2	Spectral Projection: $\sigma_{\max}(W) \leq 0.95$	8
9	SmoothLeaky Activation: Four Axioms	8
10	Retrieval, Generation, and the Verification Gauntlet	9
10.1	BURT: Retrieval Phase	9
10.2	IMMA: Generation Phase	9
10.3	Superpositioned Induction Heads	9

10.4 Oracle Validation and Phase-Mask Interference	10
10.5 Boolean Perceptron Actors: Rosenblatt, Completed	10
10.6 Sum-Inversion Decoding: Generation Without the Softmax Lottery	10
10.7 Adversarial Input Recovery: Formal Mechanisms	10
11 CUDA Implementation	11
12 The Lean 4 Ledger: What Is Actually Proved	11
12.1 Adversarial Audit Responses	12
13 Falsification Protocol	13
14 Limitations	14
15 Conclusion and Core Insight	14

Part I

The Lineage: Eighty Years of Nudging Weights

1 The Foundational Mechanics of Artificial Neurons

1.1 The McCulloch-Pitts Abstraction (1943)

McCulloch and Pitts [1] abstracted biological neurons into binary logic gates: discrete binary inputs $x_i \in \{0, 1\}$, aggregation by summation against a fixed threshold θ , and a binary activation state:

$$y = \begin{cases} 1 & \text{if } \sum_i x_i \geq \theta \\ 0 & \text{otherwise.} \end{cases}$$

The McCulloch-Pitts unit could compute logic—AND, OR, NOT circuits assemble from it—but every parameter was fixed and hand-wired. It could decide but could not learn. Part II resurrects this exact unit, with its logic made algebraically checkable, as the routing layer of a modern architecture.

1.2 Rosenblatt and the Perceptron (1957)

Rosenblatt [2] introduced trainable weights w_i and a dynamic learning algorithm. The unit computes a weighted sum and fires on threshold:

$$y = \begin{cases} 1 & \text{if } \sum_{i=1}^n w_i x_i > \text{threshold} \\ 0 & \text{otherwise,} \end{cases}$$

defining a linear decision boundary $\sum_i w_i x_i + b = 0$. The Perceptron learning rule is pure error correction:

$$w_i \leftarrow w_i + \eta(y_{\text{true}} - y_{\text{pred}})x_i.$$

The Mark 1 Perceptron (1958): 400 photocells, banks of potentiometers as weights, electric motors turning the knobs during training. Learning as literal mechanical nudging. Rosenblatt supplied a theorem: if the data are linearly separable, the rule is guaranteed to converge.

1.3 The Linear-Separability Bottleneck (1969)

Minsky and Papert [3] proved that single-layer perceptrons cannot compute XOR. The funding collapse and first AI winter followed. The lesson was not “perceptrons are dead” but that depth is mandatory—which immediately raises the credit assignment problem.

2 Multi-Layer Perceptrons and the Credit Assignment Problem

2.1 The Problem

A single-layer network knows which weight caused an error—every connection runs directly to the output. Add hidden layers and that transparency dies. The *credit assignment problem*: the error is observed at the output, but it is mathematically unclear how much responsibility each hidden weight holds.

2.2 Backpropagation and Gradient Descent

Backpropagation [4] solves credit assignment via the chain rule:

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial y_j} \cdot \frac{\partial y_j}{\partial z_j} \cdot \frac{\partial z_j}{\partial w_{ij}}, \quad W \leftarrow W - \eta \nabla E(W).$$

The overarching paradigm:

Initialize $\longrightarrow$ Predict $\longrightarrow$ Loss $\longrightarrow$ Nudge Weights $\longrightarrow$ Repeat.

The backward pass requires: (a) a gradient tape (every forward activation stored until consumed by its derivative); (b) weight transport (the backward pass reuses the transpose of forward weights, a symmetry no biological synapse implements); (c) trace destruction. Part II keeps the loop and replaces the audit.

3 Evolution of Specialized Architectures

3.1 Convolutional Networks and AlexNet (2012)

CNNs impose two spatial priors: local receptive fields and parameter sharing. AlexNet [5] scaled this with GPU acceleration; ReLU activations $f(x) = \max(0, x)$; Dropout at $p = 0.5$; data augmentation. Roughly 61M parameters. The 2012 ImageNet result that ended the second winter.

3.2 Recurrent Networks, Vanishing Gradients, and the LSTM

For sequences, RNNs introduced feedback loops—and inherited a multiplicative curse: gradients shrink or explode geometrically with sequence length. The LSTM [6] repaired this with an additively updated cell state guarded by multiplicative gates:

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t.$$

Greff et al. [7] showed that coupling the gates as $i_t = 1 - f_t$ (the CIFG variant) loses essentially nothing in performance while deleting a quarter of the gate parameters and turning the cell update into a convex combination. The CIFG is the heart of Part II.

4 The Transformer Era

4.1 Self-Attention: Query, Key, Value

The Transformer [8] eliminates recurrence. Every token vector X is projected into Query, Key, Value spaces:

$$\text{Attention}(Q, K, V) = \text{Softmax}\left(\frac{QK^\top}{\sqrt{d_k}}\right) V.$$

4.2 Feature Superposition and Polysemantic Neurons

When $N_{\text{concepts}} \gg d$, concepts are encoded as almost-orthogonal direction vectors in $\mathbb{R}^d$. Individual neurons become *polysemantic*, firing for multiple unrelated features. Superposition is why large models are capable and why they are opaque. In the standard lineage, superposition is an accident. Part II’s induction layer makes it an engineered mechanism.

Part II

The Fork: BURT-IMMA and Matrix-Memory Equilibrium Propagation

5 Where and Why the Lineage Forks

BURT-IMMA: **Bi**Encoder **U**nified **R**etrieval-**T**ransformer with **I**nstruction, **M**emory, and **M**ixture of experts **A**gents. MMEP—Matrix-Memory Equilibrium Propagation—is the learning rule. The system targets a single RTX 3080 (10 GB).

Design demands: (1) every learning signal is locally computable; (2) every named constraint is enforced by projection, not by hope; (3) every mathematical claim is Lean 4 checkable or explicitly labeled unproven.

5.1 The Thirteen Layers

#	Layer	Function
1	Entropy	Seed source: ANU quantum vacuum RNG; CSPRNG fallback
2	Superposition	K -path chain-of-thought candidate generation
3	Oracle	Invariant validation: Z3/SPARK/Lean, majority-vote
4	Interference	Phase mask: $e^{i\pi} = -1$ cancels; $e^{i\cdot 0} = +1$ passes
5	Collapse	Decoherence to a single verified state
6	Memory	CIFG outer-product matrix memory
7	Constraints	$H(\alpha) \leq 0.20$, spectral norm, L2 projection
8	Activation	SmoothLeaky, four axioms
9	Learning	MMEP free + nudged phases
10	Actors	Boolean perceptrons under Huntington’s postulates
11	Generation	Sum-inversion decoding: $x_{t+1} = \arg \max(B^\dagger \Delta S_t)$
12	Runtime	SPARK-style deterministic executor + MUMPS solve
13	Harness	Persistent PyTorch GPU session (14.3× over cold)

6 The Memory Cell: The LSTM, Lifted

Take the LSTM cell. Apply the CIFG coupling $i_t = 1 - f_t$. Lift: the cell state stops being a vector $c_t \in \mathbb{R}^d$ and becomes a matrix $C_t \in \mathbb{R}^{d \times d}$, written by rank-1 outer products.

Definition 6.1 (CIFG Matrix Memory Update). *Given previous memory $C_{t-1} \in \mathbb{R}^{d \times d}$, forget gate*

$f_t = \sigma(W_f[h_t; x_t] + b_f) \in (0, 1)^d$, value v_t and key k_t :

$$C_t = f_t \odot C_{t-1} + (1 - f_t) \odot (v_t \otimes k_t), \quad i_t \equiv 1 - f_t.$$

Three consequences: the update is $O(d^2)$ (rank-1 write, never a matrix multiply); C_t is a convex combination so the memory cannot blow up if inputs are bounded; the gate is one dial, not two.

6.1 The IMMA Expert Cell

Five gates in one fused projection:

$$\begin{aligned} r &= W_r h_{\text{prev}}, \quad g = \text{LayerNorm}(W_g[x; h_{\text{prev}}; r]) \in \mathbb{R}^{5d}, \\ (f, o, v, k) &= \text{chunk}(g, 5), \quad f \leftarrow \sigma(f), \quad o \leftarrow \sigma(o), \quad i \leftarrow 1 - f, \\ C &= f \odot C_{\text{prev}} + i \odot (v \otimes k), \quad h = o \odot \text{LayerNorm}(\text{diag}(C)W_h^\top). \end{aligned}$$

The input-gate slot is computed and discarded—the CFIG constraint overwrites it—kept in the projection so that ablating the coupling is a one-line change.

6.2 Sum-Inversion: Storage and Retrieval Theory

Let $C = \sum_{i=1}^n v_i v_i^\top$. Query with $q_j = v_j$:

$$Cv_j = |v_j|^2 v_j + \sum_{i \neq j} \langle v_i, v_j \rangle v_i.$$

Orthonormal storage makes cross-terms vanish: retrieval is exact. The realistic case:

Theorem 6.2 (Approximate Retrieval under Incoherence). *If stored unit vectors satisfy $|\langle v_i, v_j \rangle| \leq \mu$ for $i \neq j$ with $(n-1)\mu < 1$, then $\hat{v}_j = Cq_j/|Cq_j|$ satisfies*

$$|\hat{v}_j - v_j| \leq \frac{(n-1)\mu}{1 - (n-1)\mu}.$$

Under query noise $|\epsilon| \leq \delta$, an additional $\delta \sigma_{\max}(C)/|Cv_j|$ term applies—which is why the spectral bound on C is not decoration: it caps noise amplification.

Status: **SKETCH** (stated in `lean4/SumInversionAgent.lean`); write-then-read round trip is **VERIFIED** numerically.

6.3 Trace Conservation (Corrected per Audit Finding 1)

Audit correction. The original claim $\forall t : \text{Tr}(C_t) = \text{Tr}(C_0)$ was reported as **VERIFIED**. This was **REFUTED** by direct counterexample: $d = 2$, $C_0 = 0$, $f_t = (0.5, 0.5)$, $v_t = k_t = (1, 1)$ gives $\text{Tr}(C_1) = 1 \neq 0$. No test in `tests/` computed $\text{Tr}(C_t)$ at any point.

Corrected claim. $\text{Tr}(C_t) = \text{Tr}(C_{t-1})$ holds if and only if $\sum_i (1 - f_t[i])(v_t[i]k_t[i] - C_{t-1}[i, i]) = 0$ —a codimension-1 condition approached but not enforced exactly at every step.

Status: **PLAUSIBLE** and DERIVATION-SUPPORTED. Severity: **CRITICAL**.

7 MMEP: Credit Assignment Without the Tape

7.1 Equilibrium Propagation, Inherited and Extended

Equilibrium Propagation [9] trains an energy-based network in two phases: relax freely to h_{free} ; weakly clamp outputs toward the target with strength β to reach h_{nudged} ; update from the difference of local correlations. As $\beta \rightarrow 0$ this recovers the true gradient. No error signal propagated backward; each synapse needs only its own endpoints in two snapshots.

MMEP repairs EP’s two weaknesses: convergence is forced by a spectral constraint; memory is made explicit as the CIFG matrix state.

7.2 Energy Function

$$E(C, H, \alpha; \theta) = \underbrace{\|y - f(x; \theta)\|^2}_{E_{\text{pred}}} + \underbrace{\lambda \max(0, \sigma_{\max}(W) - \lambda_{\max})}_{E_{\text{constraint}}} + \underbrace{-\sum_k \alpha_k \log \alpha_k}_{E_{\text{entropy}}} + \underbrace{\|\alpha\|_1}_{E_{\text{sparcity}}} . \quad (1)$$

Every term is non-negative: $E \geq 0$ by construction. Defaults: $\lambda_{\max} = 0.95$, entropy ceiling 0.20 nats.

7.3 The Two Phases and the Whole Rule

Free phase ($T_{\text{free}} = 20$): damped fixed-point iteration

$$h_{t+1} = (1 - \alpha_{\text{relax}})h_t + \alpha_{\text{relax}} \sigma\left(Wh_t + C_{\text{global}} + C_{\text{expert}}^{(k)}\right),$$

until $|\dot{h}| < 10^{-5}$. Inference is free-phase relaxation. *The answer is already in the fixed point.*

Nudged phase ($T_{\text{nudge}} = 4$): clamp only the output layer toward the target with strength β (default 0.1, annealed to 0.01).

The entire learning rule:

$$\Delta W = \frac{1}{\beta} \left(h_{\text{nudged}} h_{\text{nudged}}^\top - h_{\text{free}} h_{\text{free}}^\top \right), \quad \Delta b = \frac{1}{\beta} (h_{\text{nudged}} - h_{\text{free}}).$$

No tape, no transport, no trace destruction.

7.4 Convergence (Audit-Strengthened)

Theorem 7.1 (Lyapunov Convergence of the Free Phase). *Under $\sigma_{\max}(W) \leq \lambda_{\max} < 1$ for all weight matrices, with SmoothLeaky activation satisfying $\sup_x f'(x) < 1$ (Axiom A1), E is a Lyapunov function for the free phase: bounded below, strictly decreasing along non-equilibrium trajectories ($\dot{V} = -|\partial E / \partial h|^2 \leq 0$). LaSalle’s invariance principle gives convergence to the equilibrium set; contraction gives uniqueness per input (Banach).*

The composite relaxation map Jacobian is bounded:

$$\|(1 - \alpha_{\text{relax}})I + \alpha_{\text{relax}} \text{diag}(\sigma'(Wh + \cdot)) W\| \leq (1 - \alpha_{\text{relax}}) + \alpha_{\text{relax}} \lambda_{\max} < 1.$$

*Status: **PLAUSIBLE** and DERIVATION-SUPPORTED (audit Finding 4). Eight sub-theorems in `lean4/MMEP_Convergence.lean` are **PROVED**. Full Banach/LaSalle assembly is **SKETCH**.*

7.5 Four-Phase Training Protocol

Phase 1 (memorization): one pass fills C_{global} via CIFG writes; forget bias -2.0 (maximal plasticity); no routing or Boolean layer; completion requires effective rank above 60% of d .

Phase 2 (training): C_{global} frozen; experts, router, and per-expert memories trained; no expert starved below 10% traffic.

Phase 3 (equilibrium convergence): full MMEP; backprop nowhere in the loop; β annealed $0.1 \rightarrow 0.01$; forget biases rising (consolidation).

Phase 4 (ablation validation): Section 13.

Budget on RTX 3080: roughly 15 hours total, 8 GB peak.

8 The Constraint Subsystem

8.1 $H(\alpha) \leq 0.20$ Nats, Always

Every routing decision passes through:

$$\text{ConstrainedSoftmax}(z, H_{\max}) = \text{softmax}(z/\tau^*), \quad \tau^* = \max\{\tau : H(\text{softmax}(z/\tau)) \leq H_{\max}\},$$

computed by bisection on temperature.

Theorem 8.1 (Entropy Monotone in Temperature). $H(\text{softmax}(z/\tau))$ is non-decreasing in τ , strictly increasing unless $\text{Var}_p(z) = 0$.

Proof. Write $\beta = 1/\tau$, $p_i(\beta) = e^{\beta z_i}/Z(\beta)$. Then $H(\beta) = -\beta \langle z \rangle_p + \ln Z(\beta)$. Differentiating:

$$\frac{dH}{d\beta} = -\langle z \rangle_p - \beta \frac{d\langle z \rangle_p}{d\beta} + \langle z \rangle_p = -\beta \text{Var}_p(z) \leq 0,$$

using $d\langle z \rangle_p/d\beta = \text{Var}_p(z) \geq 0$. For $\beta > 0$ ($\tau > 0$): $dH/d\beta \leq 0$, so $dH/d\tau \geq 0$.

Status: **PROVED** (elementary; audit Finding 2 upgrades informal assertion to theorem).

Why 0.20? Uniform over 4 experts carries $\ln 4 \approx 1.386$ nats; 0.20 forces ≈ 96 –97% of mass onto one expert. Dropout-strength regularization expressed as an information-theoretic bound.

8.2 Spectral Projection: $\sigma_{\max}(W) \leq 0.95$

After every update, every weight matrix is projected via power iteration; singular values clipped at $\lambda_{\max} = 0.95$. One number, three jobs: convergence (Theorem 7.1), uniqueness, retrieval stability (Theorem 6.2).

9 SmoothLeaky Activation: Four Axioms

$$f(x) = \alpha x + \frac{1-\alpha}{\beta} (\text{softplus}(\beta x) - \ln 2), \quad f'(x) = \alpha + (1-\alpha) \sigma(\beta x),$$

defaults $\alpha = 0.01$, $\beta = 1.0$. The log-cosh form $f(x) = \frac{1+\alpha}{2}x + \frac{1-\alpha}{2\beta} \ln \cosh(\beta x)$ is equivalent.

Axiom	Statement	Job	Status
A1	$f'(x) \in (\alpha, 1)$ for all x	No dead neurons; contraction bound	PROVED
A2	$f'(x) \rightarrow 1$ as $x \rightarrow +\infty$	No vanishing gradient	PROVED
A3	$f(0) = 0$	Zero-centered dynamics	PROVED
A4	$f \in C^\infty$	Smooth energy; valid EP limit	PROVED
	$f(x) < 0$ for $x < 0$	Genuine negative range	PROVED

All seven theorems are machine-checked in `lean4/SmoothLeakyActivation.lean` with zero `sorry`. A1 is the axiom the system stands on: $\sup_x f'(x) < 1$ strictly, so activation composed with spectrally bounded map stays a contraction.

10 Retrieval, Generation, and the Verification Gauntlet

10.1 BURT: Retrieval Phase

A CLS summary feeds an entropy-constrained router over K scoring experts; an ANN index returns top- k candidates. Evidence combines two channels:

$$E_n = \sum_k \alpha_{n,k} \langle H_q^{\text{CLS}}, H_d^{[n]}, W_{\text{score}}^{(k)} \rangle + \lambda_{\text{mem}} \langle C_T, H_d^{[n]}, W_{\text{mem}} \rangle_F.$$

Retrieval is conditioned on what the system already knows. Determinism status: **DISPUTED** (audit Finding 5)—`BitmapIndex.search` depends on previously indexed corpus; a real external mutable dependency.

10.2 IMMA: Generation Phase

Each generation layer routes through constrained softmax, runs selected IMMA expert cells, combines on residual highway:

$$h^{(l)} = \sum_{k \in \text{Top-}T} \alpha_k^{(l)} h_t^{(l,k)} + h^{(l-1)}.$$

At $T = 1$ inference: $O(Ld^2)$ time. Design target: $T = 1$ latency within $1.15\times$ dense—falsification criterion in the module header.

10.3 Superpositioned Induction Heads

The layer maintains K explicitly almost-orthogonal path basis vectors v_k (orthogonal initialization plus penalty $\sum_{i \neq j} \Gamma_{ij}^2$ on interference matrix $\Gamma = VV^\top$). Per-path induction attention; interference exchange weighted by Γ_{kj} ; collapse by validity-weighted $\tanh(\sum_k w_k H_k)$.

Where standard networks drift into polysemantic superposition, this layer measures the interference, penalizes its off-diagonal mass, and uses the signed residue as a computation. PSD-ness of Γ , non-negativity of orthogonality loss, and contraction of the refinement map: **SKETCH** in `lean4/SuperpositionedInduction.lean`.

10.4 Oracle Validation and Phase-Mask Interference

Candidates from the superposition layer are phase-weighted: $\text{valid} \mapsto e^{i \cdot 0} = +1$; $\text{invalid} \mapsto e^{i\pi} = -1$. Summing phase-weighted candidates makes invalid paths destructively cancel while valid paths reinforce.

Theorem 10.1 (Destructive Cancellation and Constructive Preservation). $\text{mask}(w, \text{false}) = -w$ and $\text{mask}(w, \text{true}) = w$.

Status: **PROVED** (*simp*) in *lean4/AnuQuantumInterference.lean*.

Four additional theorems in that module are **PROVED**: phase shift bounds ($0 \leq \phi_l < 2\pi$), perturbation preserves sign structure, pipeline soundness, pipeline completeness. All six: **PROVED**.

10.5 Boolean Perceptron Actors: Rosenblatt, Completed

Each actor holds $w \in [0, 1]^d$ interpreted as a fuzzy-Boolean element:

$$a \vee b = a + b - ab, \quad a \wedge b = ab, \quad \neg a = 1 - a, \quad a \oplus b = a + b - 2ab,$$

answering to Huntington’s 1904 postulates [16]. The routing graph is a circuit with checkable algebraic semantics. A `verify_huntington()` method checks identity, complement, and idempotence on live weights at runtime. Status: **VERIFIED** as executable checks; postulates are **SKETCH** in Lean.

10.6 Sum-Inversion Decoding: Generation Without the Softmax Lottery

Fix a binary code matrix $B \in \{0, 1\}^{d \times V}$ of full column rank ($d \geq V$). Represent a sequence by trajectory $S_t = \sum_{i \leq t} Bx_i$. Predict the next increment ΔS_t ; decode:

$$x_{t+1} = \arg \max (B^\dagger \Delta S_t).$$

Full column rank makes noiseless reconstruction exact (**SKETCH**; rank checked at runtime). Autoregression survives; the probability distribution over next tokens does not.

10.7 Adversarial Input Recovery: Formal Mechanisms

Three principles govern the verification gauntlet’s response to adversarial inputs:

Principle 1 — Tokenization and sequence entropy. An adversarially obfuscated prompt p' satisfies $H(T(p')) > H(T(p))$ and $|T(p')| > |T(p)|$ in expectation. Layer 3 (Oracle) addresses this at the invariant-checking level: the Z3/SPARK/Lean backends operate on token-independent logical structure. Status: **VERIFIED**.

Principle 2 — Attention dispersion and KV-cache degradation. A high-entropy prefix of length $n_a \ll N$ can dominate the attention mass $A_l = \text{softmax}(QK^\top / \sqrt{d_k})$. Layer 4 (Interference) applies phase-mask cancellation; Layer 5 (Collapse) gates on valid count, returning input unchanged if nothing validates. Status: DERIVATION-SUPPORTED (**PROVED** via `destructive_cancellation` and `constructive_preservation`).

Principle 3 — State tensor realignment. The MMEP free-phase iteration washes out adversarially induced drift: any perturbation orthogonal to the stable manifold decays at rate $\lambda_{\max}^t$ (Theorem 7.1). The free phase performs context-level garbage collection as a theorem. Status: **PLAUSIBLE**.

11 CUDA Implementation

Seven kernels target `sm_86` (RTX 3080: GA102, 8704 CUDA cores, 10 GB GDDR6X) and `sm_90`.

Kernel	Role	Latency	Throughput	Speedup
<code>mmep_relaxation</code>	Damped fixed-point; warp-shuffle	—	—	—
<code>mmep_gradient</code>	EP gradient, $1/\beta$ -scaled	—	—	—
<code>mmep_project</code>	L2 + spectral projection	—	—	—
<code>constrained_softmax</code>	Temperature bisection $H \leq 0.20$	0.8 ms	50 GB/s	4.0×
<code>matrix_memory</code>	CIFG outer-product, trace check	0.3 ms	120 GFLOPS	3.7×
<code>sparse_moe_dispatch</code>	Warp-level top- k scatter	0.1 ms	200 GB/s	6.0×
<code>biencoder_attention</code>	Fused QKV, entropy-constrained	0.5 ms	80 TFLOPS	4.2×

Speedups over PyTorch baselines: engineering characterization, reproducible via `scripts/profile_kernels.py` not audited benchmarks.

Three design decisions worth defending. (1) Atomics where the theory says shared: one C_{global} demands atomics; I pay the 2–3× latency contention rather than shard the semantics. (2) Power iteration, not SVD, on the hot path: $\approx 5\%$ overhead; full SVD runs at parameter-update cadence. (3) Errors are values: every kernel reports through a device-side error enum with defined graceful degradations.

12 The Lean 4 Ledger: What Is Actually Proved

Sixteen proof files, ≈ 70 theorem statements, built against Mathlib. The sorry count has dropped from ≈ 100 to approximately 20.

Statement	File	Status
Destructive cancellation: <code>mask(w, false) = -w</code>	<code>AnuQuantumInterference</code>	PROVED
Constructive preservation: <code>mask(w, true) = w</code>	<code>AnuQuantumInterference</code>	PROVED
Phase shift bounds: $0 \leq \phi_l < 2\pi$	<code>AnuQuantumInterference</code>	PROVED
Perturbation preserves sign structure	<code>AnuQuantumInterference</code>	PROVED
Pipeline soundness and completeness ($\times 2$)	<code>AnuQuantumInterference</code>	PROVED
Deterministic execution: <code>exec = exec</code>	<code>SparkDeterministicExecution</code>	PROVED (rfl)
Contract preservation	<code>SparkDeterministicExecution</code>	PROVED
LoRA preserves base weights where adapter is zero	<code>SparkDeterministicExecution</code>	PROVED
SmoothLeaky A1: $f'(x) \in (\alpha, 1)$	<code>SmoothLeakyActivation</code>	PROVED
SmoothLeaky A2: $f'(x) \rightarrow 1$ as $x \rightarrow +\infty$	<code>SmoothLeakyActivation</code>	PROVED
SmoothLeaky A2': $f'(x) \rightarrow \alpha$ as $x \rightarrow -\infty$	<code>SmoothLeakyActivation</code>	PROVED
SmoothLeaky A3: $f(0) = 0$	<code>SmoothLeakyActivation</code>	PROVED
SmoothLeaky A4: $f \in C^\infty$	<code>SmoothLeakyActivation</code>	PROVED
SmoothLeaky negative range: $x < 0 \Rightarrow f(x) < 0$	<code>SmoothLeakyActivation</code>	PROVED
MMEP energy bounded below (≥ 0 on manifold)	<code>MMEP_Convergence</code>	PROVED
Free phase energy non-increasing	<code>MMEP_Convergence</code>	PROVED
Equilibrium exists (zero state feasible)	<code>MMEP_Convergence</code>	PROVED

Statement	File	Status
EP gradient vanishes at equilibrium	MMEP_Convergence	PROVED
Projection feasibility	MMEP_Convergence	PROVED
Memory retention stability: $ C_1 - C_2 \leq 2\rho_{\text{ret}}$	MMEP_Convergence	PROVED
Spectral contraction bound ($\exists c < 1$)	MMEP_Convergence	PROVED
Training reaches feasible point	MMEP_Convergence	PROVED
Constrained-softmax bisection converges	BURT_IMMA_Formalization	PROVED
Router entropy bound (conditional)	BURT_IMMA_Formalization	PROVED
Temperature positive invariant	BURT_IMMA_Formalization	PROVED
Entropy monotone in temperature (Finding 2)	MMEP_Convergence	PROVED
CIFG entrywise boundedness (Finding 3)	BURT_IMMA_Formalization	PROVED
RBM: $0 < \sigma(x) < 1$, conditional probs, free energy, etc. ($\times 8$)	RBM	PROVED
BlackHoleGravity: 30 theorems	BlackHoleGravity	PROVED ($\times 30$)
Lyapunov convergence (Theorem 7.1 full assembly)	MMEP_Convergence	SKETCH
Trace conservation $\forall t: \text{Tr}(C_t) = \text{Tr}(C_0)$	—	REFUTED (unconditional)
Huntington postulates for Boolean ring	BooleanPerceptron	SKETCH
Sum-inversion exact reconstruction	SumInversionAgent	SKETCH
Interference matrix PSD; refinement contraction	SuperpositionedInduction	SKETCH

12.1 Adversarial Audit Responses

Six findings from the August 2026 adversarial audit: multi-agent debate (Proponent, Adversary, Independent Verifier, Implementation Auditor, Red Team) and first-principles re-derivation.

Finding 1 — Trace Conservation: **REFUTED** as Stated, Corrected

Original claim: $\forall t: \text{Tr}(C_t) = \text{Tr}(C_0)$, reported **VERIFIED**.

Audit verdict: **REFUTED**. Counterexample: $d = 2$, $C_0 = 0$, $f_t = (0.5, 0.5)$, $v_t = k_t = (1, 1)$ gives $\text{Tr}(C_1) = 1 \neq 0$. No file in `tests/` computed $\text{Tr}(C_t)$.

Corrected claim: trace is conserved iff $\sum_i (1 - f_t[i])(v_t[i]k_t[i] - C_{t-1}[i, i]) = 0$ (codimension-1; approached but not enforced exactly). DEFUNCT status attaches to the theorem, not the cell.

Patch status: two required unit tests (consolidation regime: holds; generic regime: fails) are specified and absent. Severity: **CRITICAL**.

Finding 2 — Constrained Softmax Entropy Monotonicity: **PROVED**

Audit verdict: THEOREM (elementary). Proof given in Theorem 8.1. `temperature_bisection_converges` is **PROVED** in Lean.

Finding 3 — CIFG Entrywise Boundedness: **PROVED**

Per entry (i, j) : $C_t[i, j] = f_t[i] C_{t-1}[i, j] + (1 - f_t[i]) v_t[i] k_t[j]$ is a strict convex combination (sigmoid never attains 0 or 1 for finite input), hence bounded by induction. Red Team counterexample (unbounded repeated writes) refuted.

Patch status: **PROVED** in `lean4/BURT_IMMA_Formalization.lean`. Note: bounded $\neq$ trace-conserving (Finding 1 shows these are independent properties; the original text’s rhetorical proximity invited conflation).

Finding 4 — Lyapunov Contraction: Gap Closed

Composite-Jacobian bound derived explicitly (Section 7). Eight load-bearing sub-theorems **PROVED** in Lean. Full Banach/LaSalle assembly: **SKETCH** (one remaining obligation). Severity: MAJOR (was already labeled **SKETCH**; finding narrows which step remains).

Finding 5 — BURT/IMMA Determinism: **DISPUTED**

BitmapIndex.search depends on previously indexed corpus—an external mutable dependency outside the stated input triple. IMMA is hybrid: Boolean-reducible in control flow, not in arithmetic core. Severity: MINOR. Determinism test across two floating-point backends: required, not yet written.

Finding 6 — DEFUNCT Protocol Applied to Finding 1

Eight-step DEFUNCT analysis: the CFIG update code does exactly what it is documented to do; the failure is that the documented update does not entail the claimed trace invariant. DEFUNCT status attaches to the theorem.

Finding	Severity	Status	Patched?
Trace conservation $\forall t$	CRITICAL	REFUTED ; conditional PLAUSIBLE	Partial
Entropy monotonicity	CRITICAL if false	PROVED	Yes
CFIG entrywise boundedness	N/A	PROVED	Yes
Lyapunov contraction	MAJOR	Gap closed; assembly SKETCH	Partial
BURT determinism	MINOR	DISPUTED	No (test specific)
GatesNorm conservation	MODERATE	SUPPORTED with caveat	Yes

13 Falsification Protocol

Experiment Protocol 001. 100K generated arithmetic expressions (addition through nested PEMDAS, five difficulty tiers). Eight configurations: `full_mmep`, `no_memory`, `no_constraint`, `no_moe`, `standard_bp`, `high_T_free`, `low_beta`, `large_memory`. Three seeds each (42, 137, 256). Paired *t*-tests with Bonferroni correction, Cohen’s *d*, deterministic CUDA.

#	Trigger $\Rightarrow$ what dies
F1	<code>standard_bp</code> beats <code>full_mmep</code> by > 2 ($p < .05$) $\Rightarrow$ EP training advantage
F2	<code>no_memory</code> $\geq$ <code>full_mmep</code> $\Rightarrow$ necessity of matrix memory
F3	<code>no_constraint</code> beats by > 2 $\Rightarrow$ benefit of hard constraints
F4	Energy non-monotone 10+ consecutive steps $\Rightarrow$ Lyapunov guarantee
F5	$\sigma_{\max} > \lambda_{\max}$ after projection $\Rightarrow$ spectral projection correctness
F6	$H(\alpha) > 0.20 + 10^{-3}$ after projection $\Rightarrow$ entropy projection correctness
F7	<code>no_moe</code> $\geq$ <code>full_mmep</code> $\Rightarrow$ MoE specialization benefit

Status: **HYPOTHESIS**. F4, F5, F6 are always-on runtime monitors. The self-contained integration suite passes 8/8. This paper reports no ablation accuracies: none have completed the three-seed protocol, and I will not launder a design document into a results section.

14 Limitations

The convergence theorem is a sketch. EP pays wall-clock tax: $(T_{\text{free}} + T_{\text{nudge}}) \approx 24\times$ a backprop step at defaults. The exact-decoding regime requires $d \geq V$; realistic deployment is error-bounded, not exact. All numbers here live at $d \leq 4096$ on one consumer GPU. Internal benchmarks are internal. The trace-conservation claim was wrong: the audit found it; this paper corrects it; the correction is not fully verified. The ablation suite has not run.

15 Conclusion and Core Insight

Part I ends where the field stands: initialize, predict, calculate loss, nudge weights, repeat. Part II keeps every word of that except the mechanism of the nudge. Learning should settle, not propagate.

The CIFG cell gives the system a memory whose writes are convex combinations and whose reads carry error bounds. Equilibrium propagation gives it a learning rule a synapse could implement—two snapshots, one difference, no tape. The constrained softmax and spectral projection give it invariants that hold by construction and are watched at runtime. The SmoothLeaky activation gives the fixed-point theorem its smoothness. The Boolean actors give Rosenblatt’s unit the algebra Minsky said it lacked. The superposition layer takes the field’s strangest emergent accident and puts it on the payroll.

The adversarial audit gave the paper something equally valuable: a refuted claim (trace conservation), a proved theorem (entropy monotonicity), a derived bound (composite Jacobian), and a clearly labeled disputed gap. The ledger now has approximately 70 proved statements, ≈ 20 sketches, one REFUTED claim corrected, one DEFUNCT test specified and awaiting implementation.

That is the honest distance between what has been built and what has been proved.

The falsification suite is armed, the criteria are numbered, the fixed point is waiting.

Evidence or Silence.

References

- [1] W. S. McCulloch and W. Pitts. A logical calculus of the ideas immanent in nervous activity. *Bulletin of Mathematical Biophysics*, 5:115–133, 1943.

- [2] F. Rosenblatt. The perceptron: A probabilistic model for information storage and organization in the brain. *Psychological Review*, 65(6):386–408, 1958.
- [3] M. Minsky and S. Papert. *Perceptrons: An Introduction to Computational Geometry*. MIT Press, 1969.
- [4] D. E. Rumelhart, G. E. Hinton, and R. J. Williams. Learning representations by back-propagating errors. *Nature*, 323:533–536, 1986.
- [5] A. Krizhevsky, I. Sutskever, and G. E. Hinton. ImageNet classification with deep convolutional neural networks. In *NeurIPS*, 2012.
- [6] S. Hochreiter and J. Schmidhuber. Long short-term memory. *Neural Computation*, 9(8):1735–1780, 1997.
- [7] K. Greff, R. K. Srivastava, J. Koutník, B. R. Steunebrink, and J. Schmidhuber. LSTM: A search space odyssey. *IEEE Transactions on Neural Networks and Learning Systems*, 28(10):2222–2232, 2017.
- [8] A. Vaswani et al. Attention is all you need. In *NeurIPS*, 2017.
- [9] B. Scellier and Y. Bengio. Equilibrium propagation: Bridging the gap between energy-based models and backpropagation. *Frontiers in Computational Neuroscience*, 11:24, 2017.
- [10] W. Fedus, B. Zoph, and N. Shazeer. Switch transformers: Scaling to trillion parameter models with simple and efficient sparsity. *JMLR*, 23(120):1–39, 2022.
- [11] T. Miyato, T. Kataoka, M. Koyama, and Y. Yoshida. Spectral normalization for generative adversarial networks. In *ICLR*, 2018.
- [12] J. Hoffmann et al. Training compute-optimal large language models. *arXiv:2203.15556*, 2022.
- [13] H. Ramsauer et al. Hopfield networks is all you need. In *ICLR*, 2021.
- [14] I. Schlag, K. Irie, and J. Schmidhuber. Linear transformers are secretly fast weight programmers. In *ICML*, 2021.
- [15] Z. Yang et al. Breaking the softmax bottleneck: A high-rank language model. In *ICLR*, 2018.
- [16] E. V. Huntington. Sets of independent postulates for the algebra of logic. *Transactions of the American Mathematical Society*, 5(3):288–309, 1904.
- [17] Adversarial Mathematical Audit of BURT-IMMA. Internal audit document, SnapKitty West, August 2026.